



Sizeof and Vtables; Modern C++ Topics

ITP 435
Week 3, Lecture 1



Sizeof and Vtables

sizeof Operator



Tells you the number of bytes a particular data type (or variable) takes up:

```
// Returns 4  
sizeof(int)
```

```
char a;  
// Returns 1  
sizeof(a)
```

```
// Returns 1  
sizeof(bool)
```

sizeof Question #1



```
int* ptr;  
ptr = new int[10];
```

Q: What is the `sizeof(ptr)`??

A: On a 32-bit system, 4. On a 64-bit system, 8 (we'll stick with 64-bit).

sizeof Question #2



```
int array[10];
```

Q: What is the `sizeof(array)`??

A: 40, because each `int` is 4 bytes in size.

sizeof Question #3



```
int array[10];  
int* ptr = array;
```

Q: What is the `sizeof(ptr)`??

A: It's a pointer, so you'll get either 4 or 8 again!

sizeof Question #4



```
class Test1
{
    char c;
    int i;
};
```

Q: What is the `sizeof`(Test1) ??

A: 8, because of padding.

Class Padding



- By default, Visual Studio (and most compilers) guarantees alignment to be equal to the size of the largest basic type in the struct/class
- ***x-byte aligned*** means the variable's memory address is guaranteed to be divisible by x
- So our previous example, we have a 4 byte variable (`int i`) that must be 4-byte aligned. Thus, we need padding:

| Test1 | |
|-------|---------|
| c | 3 BYTES |
| i | |

More Padding



What happens with this?

```
class Test2
{
    char a;
    int i;
    char b;
};
```

| Test2 | |
|-------|---------|
| a | 3 BYTES |
| i | |
| b | 3 BYTES |

`sizeof(Test2) == 12`

The compiler won't rearrange variables!!

Even More Padding



Solution: Always declare your variables from largest to smallest:

```
class Test3
{
    int i;
    char a;
    char b;
};
```

| Test3 | | |
|-------|---|---------|
| i | | |
| a | b | 2 bytes |

sizeof Question #5



Q: What's the sizeof this class?

```
class VirtualClass
{
    int i;
    virtual void F();
};
```

A: 16 (for a 64-bit program). Because it has a virtual function, the size of the class is the pointer to a virtual function table + the int+padding.

sizeof Question #5



Q: What's the sizeof this class?

```
class VirtualClass
{
    int i;
    virtual void F();
};
```

| VirtualClass | |
|--------------|---------|
| vtable* | |
| i | 4 bytes |

A: 16 (for a 64-bit program). Because it has a virtual function, the size of the class is the pointer to a virtual function table + the int+ padding.



Memory cost:

- Need to add 1 pointer to the start of a class' data once it has virtual functions.
- This points to the “virtual function table” or vtable.

Performance cost:

- At run-time, when a virtual function is called virtually, the vtable pointer is dereferenced, which then finds the correct function to call from the vtable.

Padding and Aggregate Types



```
struct A {  
    double d;  
    char s;  
};  
  
struct X {  
    virtual void a();  
    virtual void b();  
    void c();  
    virtual ~X();  
    A myA;  
};
```

| X | |
|-------------|-----------------|
| vtable* (8) | |
| myA.d (8) | |
| myA.c (1) | 7 bytes padding |

- Even though the `sizeof(A)` is 16, the “row size” is still 8 because we only care about the largest basic type inside the aggregate or in the current class

Virtual Function Tables



- Simple example:

```
class A
{
public:
    virtual void x();
    virtual void y();
    virtual ~A();
};
```

```
class B : public A
{
public:
    virtual void x();
    virtual void z();
    virtual ~B();
};
```

| vtable for A (24 bytes if 64-bit) | |
|-----------------------------------|------------------|
| Index | Function Pointer |
| 0 | &A::x() |
| 1 | &A::y() |
| 2 | &A::~A() |

| vtable for B (32 bytes if 64-bit) | |
|-----------------------------------|------------------|
| Index | Function Pointer |
| 0 | &B::x() |
| 1 | &A::y() |
| 2 | &B::~B() |
| 3 | &B::z() |

A couple more sizeof/vtable examples...



```
class X
{
    int i;
public:
    virtual void a();
    void b();
    virtual void c();
    virtual ~X();
};
```

`sizeof(X) == 16`

| X | |
|---------|---------|
| vtable* | |
| i | padding |


| vtable for X (24 bytes if 64-bit) | |
|-----------------------------------|------------------|
| Index | Function Pointer |
| 0 | &X::a() |
| 1 | &X::c() |
| 2 | &X::~~X() |

A couple more sizeof/vtable examples...



```
class Y : public X
{
    char p;
public:
    virtual void c() override;
};
```

`sizeof(Y) == 24`



| vtable for Y (24 bytes if 64-bit) | |
|-----------------------------------|------------------|
| Index | Function Pointer |
| 0 | &X::a() |
| 1 | &Y::c() |
| 2 | &X::~~X() |

In-class activity



```
class A
{
public:
    virtual void f();
    virtual void g();
    void h();
private:
    char c;
    double d;
};
```

```
class B : public A
{
public:
    void f() override;
    void h();
    virtual void i();
private:
    int i;
};
```

sizeof Question #6



Q: What's the sizeof this class?

```
class Empty  
{  
};
```

A: 1. Classes cannot take up 0 memory.



- A **union** is a struct that can hold only one of its members at any time*
- Example:

```
union MyUnion
```

```
{
```

```
    // Only can hold one of these at a time
```

```
    int n;
```

```
    float f;
```

```
    char c;
```

```
};
```

*Technically, the union just takes the same chunk of memory and can treat it as different types as desired

Example



```
MyUnion test;  
test.n = 50;
```

| |
|---------|
| MyUnion |
| 50 |

Example



```
MyUnion test;  
test.n = 50;  
  
test.f = 2.5f;
```

MyUnion

2.5f

Example



```
MyUnion test;
```

```
test.n = 50;
```

```
test.f = 2.5f;
```

```
test.c = 'a';
```

| |
|---------|
| MyUnion |
| 'a' |



- The sizeof a union is equal to the size of the largest member of the union:

```
union MyUnion
{
    int n;    // 4 bytes
    float f;  // 4 bytes
    char c;   // 1 byte
};
```

- sizeof(myUnion)????
- 4

Why use a union?



There's a couple of cases where it might be useful, maybe

1. Addressing the same chunk of memory in multiple ways:

```
union PointUnion
{
    struct { float x, y, z; } p;
    float a[3];
};
```

```
// Later...
```

```
PointUnion pt;
pt.p.x = 5.0f;
std::cout << pt.a[0]; // 5.0f
```

Why use a union? Cont'd



2. Can use to pack in data that's mutually exclusive:

```
struct SomeFileType
{
    enum SubType { v1, v2 };
    SubType type;
    union SomeUnion
    {
        // Could be one of these, based on sub type
        struct { /* */ } s1;
        struct { /* */ } s2;
    };
};
```

Why use a union? Cont'd



- In practice, sort of an edge case
- More of a C thing than C++
- YMMV



Modern C++ Topics



- C++11 feature to allow *compile-time* computation

- Example:

```
constexpr int max(int a, int b)
{
    return (a > b ? a : b);
}
```

- Then if we have code like this:

```
constexpr int a = max(5, 6);
```

- The *compiler* will replace it with:

```
int a = 6;
```



- Better example:

```
constexpr int factorial(int x)
{
    if (x > 0)
    {
        return x * factorial(x - 1);
    }
    else
    {
        return 1;
    }
}
```

An Annoyance



- This won't compile:

```
struct Test
{
    static const int CONST_INT = 0;
    static const float CONST_FLOAT = 0.0f;
};
```

- Error: “A member of type const float cannot have an in-class initializer”

But this does!



```
struct Test
```

```
{
```

```
    static constexpr int CONST_INT = 0;
```

```
    static constexpr float CONST_FLOAT = 0.0f;
```

```
};
```




- “Item 15: Use constexpr whenever possible”
- I think this is a decent idea



- Some samples here:
- <https://github.com/chalonverse/CPPBeyondSamples>



Filesystem (Ex02 in repo)



- Purpose:
“implementations of an interface that computer programs written in the C++ programming language may use to perform operations on file systems and their components, such as paths, regular files, and directories.”
- Gives us a cross-platform way for common file system operations like finding if a file exists, getting the size of files, iterating over files in directories, and a lot more!

Checking if Files exist



```
// Shortcut so we don't have to type out std::filesystem every time
namespace fs = std::filesystem;

// Get the size of a file (throws an exception if it doesn't exist)
try
{
    auto fileSize = fs::file_size("CMakeLists.txt");
    std::cout << "Size = " << fileSize << "\n";
}
catch (fs::filesystem_error& e)
{
    std::cout << "Error from filesystem: " << e.what() << "\n";
}
```

No exceptions?



- If you don't want to use exceptions, most versions of filesystem functions have an alternative version that takes in an error_code:

```
std::error_code ec;  
auto fileSize = fs::file_size("asdf", ec);  
  
// A default error code means it's ok  
if (ec == std::error_code{}) {  
    std::cout << "Size = " << fileSize << "\n";  
} else {  
    std::cout << "Error: " << ec.message() << "\n";  
}
```



- You can construct a path object to a single file, and extract information from it:

```
fs::path pathTest("CMakeLists.txt");

// Does this file exist?
if (fs::exists(pathTest))
{
    std::cout << "File exists\n";
}
else
{
    std::cout << "File does not exist\n";
}
```

Iterating over all files in a directory



```
for (const auto& p : fs::directory_iterator(".")) {  
    std::cout << p.path() << "\n";  
}
```

Sample output:

```
".\\.git"  
".\\.gitignore"  
".\\.travis.yml"  
".\\.vs"  
".\\build"  
".\\CMakeLists.txt"  
".\\CMakeSettings.json"  
".\\Ex01"  
".\\Ex02"  
".\\Ex03"  
".\\Ex04"  
".\\Ex05"  
".\\Ex06"  
".\\Ex07"  
".\\README.md"
```




Optional, Variant, Any (Ex03 in repo)



- An optional is a wrapper that can ***optionally*** contain a value
- You can use this for cases where you're not sure whether or not something would happen, for example:

```
// This function will try to convert a string to an integer and either --  
// Return the integer if successful  
// Return an unset optional if not successful  
std::optional<int> TryGetInt(const std::string& s) {  
    try {  
        int result = std::stoi(s);  
        return std::optional<int>(result);  
    } catch (std::exception&) {  
        // stoi failed so return an unset optional  
        return std::optional<int>();  
    }  
}
```



- You can pass around optionals to other functions, and the bool conversion operator will be true if it's set, false otherwise:

```
void CoutOptional(const std::optional<int>& o) {  
    if (o) {  
        std::cout << o.value() << "\n";  
    } else {  
        std::cout << "Optional was unset!\n";  
    }  
}
```



- This code:

```
auto result1 = TryGetInt("10");  
CoutOptional(result1);  
auto result2 = TryGetInt("xyzw");  
CoutOptional(result2);
```

- Would output:

10

Optional was unset!



- This is C++'s answer to a union
- You set all the possible types as template arguments to the variant when you create it:

```
// This variant can either contain an int or a string  
std::variant<int, std::string> var;
```



- You can set a variant just like a normal variable of any of the types
- To get the variant, use the get template function:

```
// Set it to an int
```

```
var = 20;
```

```
// We can use get<type>(variant) to extract the type
```

```
std::cout << "As int: " << std::get<int>(var) << "\n";
```

```
// Set it to a string
```

```
var = "Hello";
```

```
std::cout << "As string: " << std::get<std::string>(var) << "\n";
```



- If you access the variant as a type it's not currently set to, you get an exception
- This means variants, unlike unions, are type-safe

```
// Set it to a string
```

```
var = "Hello";
```

```
std::cout << "As string: " << std::get<std::string>(var) << "\n";
```

```
// If we try to access it as a string we have an exception
```

```
try {
```

```
    int result = std::get<int>(var);
```

```
} catch (std::bad_variant_access&) {
```

```
    std::cout << "Error, variant wasn't set to an int!\n";
```

```
}
```



- Like variant, but not restricted to a specific type
- Not as efficient as using variant

```
// Initialize an any to any type
auto any1 = std::any(50);
// To extract the value we have to use any_cast
std::cout << "As int: " << std::any_cast<int>(any1) << '\n';
```




```
// If it doesn't contain the type, any_cast throws an exception
try
{
    std::cout << "As string: " <<
        std::any_cast<std::string>(any1) << '\n';
}
catch (std::bad_any_cast&)
{
    std::cout << "Any cast failed :(" << std::endl;
}
```



Feature Testing (C++20 and later)



- Need to include the `<version>` header
- Allows you to test, with the preprocessor, if a feature is supported, like:

```
#if defined(__cpp_lib_filesystem)
// This code is compiled if the platform
// support std::filesystem
std::filesystem::path myPath("hello");
#endif
```

Better way of checking for parallel algos



```
#if !defined(__cpp_lib_parallel_algorithm)
#if defined(__APPLE__)
// On Apple platforms, we can use this if we have to
#define PSTLD_HEADER_ONLY
#define PSTLD_HACK_INTO_STD
#include "pstld.h"
#else
#error "Parallel algos are required"
#endif // defined(__APPLE__)
#else
#include <algorithm>
#include <execution>
#endif // !defined(__cpp_lib_parallel_algorithm)
```