



Functional Programming; Threads

ITP 435
Week 5, Lecture 2



Lambdas Review

Lambda Expressions – Motivation



- By default, `std::sort` sorts in ascending order:

```
std::vector<int> v;  
v.emplace_back(10);  
v.emplace_back(-10);  
v.emplace_back(100);  
std::sort(v.begin(), v.end());  
// v = {-10, 10, 100}
```

- What if we want descending order?

A Custom Comparator!



```
std::sort(v.begin(), v.end(), [](int a, int b) {  
    // Return a greater than b instead of less  
    return a > b;  
});  
// v = {100, 10, -10}
```

- This uses a new syntax...



- A *lambda expression* is an inline declaration of a function and implementation

```
[](int a, int b) {  
    // Return a greater than b instead of less  
    return a > b;  
}
```

- In other languages, this may be called anonymous functions (like in JavaScript)
- Let's explain this crazy syntax!

Lambda Expression Syntax



Capture Clause
(empty in this case)

Function Parameters

Body of
function

```
[](int a, int b) {  
    return a > b;  
}
```

Lambda Expression Syntax, Another Example



Capture Clause

Function Parameters

Return Value
(optional)

```
[&count](int i) -> int  
{
```

```
    std::cout << i << std::endl;  
    count++;  
    return count;  
}
```

Body



- This is used to “capture” variables that exist outside the lambda expression, and use them inside the lambda
- Variables can be captured by value or reference:

```
// Capture ALL local variables by value (not recommended)
```

```
[=]
```

```
// Capture ALL local variables by reference (not recommended)
```

```
[&]
```

```
// Capture x by value and y by reference
```

```
[x, &y]
```

```
// Capture count by reference, and all other locals by value
```

```
[=, &count]
```

```
// Capture this by value (can't be captured by reference)
```

```
// If you want to use any member functions or variables in the
```

```
// lambda, you have to capture this.
```

```
[this]
```




Lambdas stored in variables

- You're allowed to store lambdas in a variable
- Instead of figuring out the type, use auto
- Our previous example could be:

```
auto greater = [](int a, int b) {  
    return a > b;  
};  
std::sort(v.begin(), v.end(), greater);  
// v = {100, 10, -10}
```

Specifying the Type of a Lambda



- If you need to specify the type, use the `std::function` template class (in the `<functional>` header):

```
std::function<bool(int, int)> greater =  
    [](int a, int b) {  
        return a > b;  
    };
```

- (This says the lambda returns a `bool` and takes in two `ints` as parameters)



Event Class Example

Event Class



```
class Event
{
public:
    // Add a handler for this event
    void AddHandler(std::function<void()> handler)
    {
        mHandlers.emplace_back(handler);
    }
    // Loop over handlers and call each lambda
    void Trigger()
    {
        for (auto& f : mHandlers)
        {
            f();
        }
    }
private:
    std::vector<std::function<void()>> mHandlers;
};
```

Event Class in action – What's the output?



```
std::string name = "Sanjay";
```

```
Event myEvent;
```

```
myEvent.AddHandler([]() {
```

```
    std::cout << "This is handler 1" << std::endl;
```

```
});
```

```
myEvent.AddHandler([name]() {
```

```
    std::cout << "This is handler 2: " << name << std::endl;
```

```
});
```

```
int x = 5;
```

```
myEvent.AddHandler([&x]() {
```

```
    x++;
```

```
});
```

```
std::cout << "x = " << x << std::endl;
```

```
myEvent.Trigger();
```

```
std::cout << "x = " << x << std::endl;
```

Output



```
x = 5
```

```
This is handler 1
```

```
This is handler 2: Sanjay
```

```
x = 6
```



Omitting the Return Type

- If you omit the return type, the return type is deduced from the return statements (much like how auto deduces type):

```
auto lambda = [](int x) {  
    // This will return a bool  
    return x == 5;  
};
```

```
auto lambda2 = [](int a) {  
    // This returns an int  
    if (a < 0) {  
        return a * -1;  
    } else {  
        return a;  
    }  
};
```

- Otherwise, it's assumed the lambda returns void



Explicit Return Type

- You can explicitly specify the return type using a special return syntax:

```
auto lambda = []() -> void {  
    // This returns void  
};
```

```
auto lambda2 = []() -> int {  
    // This returns an int  
    return 0;  
};
```




Functional Programming



- What is functional programming?
- Here's a great quote: "functional programming is the mustachioed hipster of programming paradigms"





- In functional programming, functions are *first-class citizens*
- This means functions can...
 - Be passed as arguments to other functions
 - Can be returned from other functions
 - Can be assigned to a variable or stored in a container
- ***C++ Lambda expressions satisfy these criteria!***



- In functional programming, *higher-order functions* – functions that can accept other functions as arguments – are allowed
- Three very common higher-order functions:
 - **map** – Apply a function to each element in a collection, storing the result in another collection
 - **filter** – Remove elements from a collection based on a filtering function
 - **reduce** – Reduce a collection to a single value by applying a binary operation repeatedly
- All three of these higher-order functions are more or less supported in C++, though they have different names



- Generally, functional programming tries to largely limit *side effects*
 - having any interaction with the world “outside” from the instance of a function call
- Examples of side effects:
 - Modifying data passed into the function (eg. pass-by-reference where you change the values)
 - Modifying global or static data
 - **Any I/O** (*can be hard to realistically enforce*)



- Functional programs should generally be stateless – in a practical sense, this means no globals, statics, etc
- All functions should accept at least one argument
- All functions must return data or another function – no void functions
- Avoid loops/iteration – always use recursion or higher-order functions. Some functional programming languages do not have any iteration at all

Average a Vector – An Example



```
// Imperative-style implementation
float averageVector(const std::vector<float>& v) {
    float sum = 0.0f;
    for (auto& i : v) {
        sum += i;
    }

    return sum / v.size();
}
```

Average a Vector, Cont'd



- To follow the principles of functional programming, the loop has to go...

```
float averageVector(const std::vector<float>& v) {  
    float sum = 0.0f;  
    for (auto& i : v) {  
        sum += i;  
    }  
  
    return sum / v.size();  
}
```




- **Functional decomposition** is a fancy way of saying “break a function into sub functions”
- Let’s make an average function that computes the average given a sum and a quantity:

```
float average(float sum, size_t qty) {  
    return sum / qty;  
}
```



- **reduce** – take a collection and reduce it to a single value by applying a binary operator repeatedly
- ***We want to reduce the collection to the sum of its components!***
- In C++, reduce can be implemented via `std::accumulate` in the `<numeric>` header

Sum Vector



```
float sumVector(const std::vector<float>& v) {  
    // std::accumulate can work as a REDUCE  
    return std::accumulate(v.begin(), // Start of range  
                           v.end(),   // End of range  
                           0.0f,      // Initial value  
                           // Binary lambda expression  
                           [](const float& a, const float& b) {  
                               return a + b;  
                           }  
    );  
}
```

Putting it Together



```
// Functional implementation
```

```
float average(float sum, size_t qty) {  
    return sum / qty;  
}
```

```
float sumVector(const std::vector<float>& v) {  
    return std::accumulate(v.begin(), v.end(), 0.0f);  
}
```

```
float averageVector(const std::vector<float>& v) {  
    return average(sumVector(v), v.size());  
}
```

all_of, any_of, none_of – A lambda usage case



- A series of new functions in C++11 in `<algorithm>`
- Given a range of values and a unary predicate with this signature:
`bool predicate(const T& a);`
- It'll return true if ***all of***, ***any of***, or ***none of*** the elements satisfy the condition

all_of Example



```
std::vector<int> v1{ 2, 4, 6, 8, 10 };  
if (std::all_of(v1.begin(), v1.end(), [](const int& i) {  
    return (i % 2) == 0;  
})))  
{  
    std::cout << "All are even!" << std::endl;  
}  
else  
{  
    std::cout << "All aren't even" << std::endl;  
}
```

copy_if – Another Example



```
std::vector<int> from{ 1, 2, 3, 4, 5 };  
std::vector<int> to;
```

```
auto is_odd = [](const int& i) { return i % 2 == 1 };
```

```
// Copy from the "from" container into the "to" container,  
// only if is_odd returns true for that element  
std::copy_if(from.begin(), from.end(),  
             std::back_inserter(to), is_odd);
```



- Copies n elements from a source collection to a back_inserted collection



- The **map** higher-order function applies a function to each element, saving the results in a different collection
- It can be approximated by the `std::transform` function:

```
std::vector<float> divEachBy(const std::vector<float>& v,
                             float denominator) {
    std::vector<float> ret;
    // std::transform can be used to map
    std::transform(v.begin(), // Start of range
                  v.end(), // End of range
                  std::back_inserter(ret), // Collection to insert into
                  // Unary Function that returns transform value
                  [denominator](const float& a) {
                      return a / denominator;
                  });
    return ret;
}
```

Write our own “map” function?



```
template <typename T, typename U>
T map(const T& v, U f) {
    T ret;
    std::transform(v.begin(), v.end(),
                  std::back_inserter(ret), f);
    return ret;
}
```

- (Aside: This is an example of where using namespace std may not be good)

Rewriting divEachBy



- Now we can simplify the code:

```
std::vector<float> divEachBy(const std::vector<float>& v,  
                             float denominator) {  
    return map(v, [denominator](const float& a) {  
        return a / denominator;  
    });  
}
```



- The *filter* higher-order function removes elements based on a Boolean filter condition
- We could write this higher-order function, too:

```
template <typename T, typename U>
T filter(const T& v, U f) {
    T ret;
    std::copy_if(v.begin(), // Start of range
                v.end(),    // End of range
                std::back_inserter(ret), // Where to insert
                f);         // Boolean unary function
    return ret;
}
```



```
std::list<int> filterIsPositive(const std::list<int>& l) {  
    return filter(l, [](const int& a) {  
        return a > 0;  
    });  
}
```





- Helper that can be used to insert elements into a collection, where elements need to be initialized with a particular function:

```
// Make a vector w/ 10 elements
std::vector<int> v(10);
std::generate(v.begin(), // start of range
              v.end(),    // end of range
              // Lambda that returns generated value
              []() {
                  return rand();
              });
```

adjacent_difference



```
std::vector<int> from({ 10, 15, 20, 25, 30, 35 });
std::vector<int> to;
std::adjacent_difference(from.begin(), // Start of range
    from.end(), // End of range
    std::back_inserter(to), // Collection to insert into
    // Lambda to compute difference (defaults to subtraction)
    [](int a, int b) {
        return a - b;
    });
```

- The first element in the resulting range is a copy, the rest of the collection is differences (index 1 – index 0 of the source is stored in index 1 of the destination, and so on)
- So in the above, “to” will contain:
{ 10, 5, 5, 5, 5, 5 }



Threads Basics

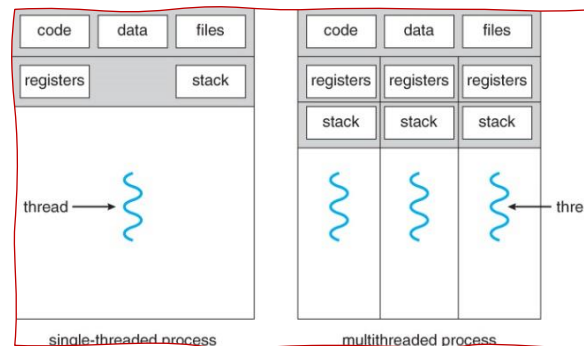



- A computer program:
 - A passive collection of instructions typically stored in a file on disk.
- A process is a program in execution
 - **Allocated memory** (heap, stack, code).
 - **OS descriptors** of resources that are allocated to the process: such as file descriptors.
 - **Security attributes**, such as the process owner and the process' set of permissions.
 - **Processor state** (context), such as the content of registers and physical memory addressing.
- Several processes may be associated with the same program
 - Multiple instances of a program running at the same time.




Thread

- The smallest sequence of instructions that can be managed independently by a OS.
- Multiple threads can exist within one process
 - Can execute concurrently and sharing resources such as memory, while different processes do not share these resources.
- Threads of a process share
 - Code
 - Memory (Global variables, heap)




Threads





What Do We Cover?

- **Threads**: Create using `std::thread`
 - Doesn't return a value
- **Tasks**: Create using `std::async`
 - Returns a value
- Both can use:
 - Pointer to function
 - Functor
 - Lambda functions





`std::thread`

Create Threads Using Function Pointers




Calculating the sum of numbers in the range [**start**, **end**)

```
void AccumulateRange(uint64_t &sum, uint64_t start,
                    uint64_t end) {
    sum = 0;
    for (uint64_t i = start; i < end; i++) {
        sum += i;
    }
}
```



Calculating the sum of numbers in the range [**start**, **end**)

Two threads each working on half of the range



partial_sum[0]

partial_sum[1]

total = **partial_sum[0]** + **partial_sum[1]**



Pointer to Function

```
void AccumulateRange(uint64_t &sum, uint64_t start,
                    uint64_t end) {
    sum = 0;
    for (uint64_t i = start; i < end; i++) {
        sum += i;
    }
}
```

Function

```
std::thread t1(AccumulateRange,
```

```
std::ref(partial_sums[0]), 0, 1000/2 );
```

Function
parameters



All parameters are passed by **Value**

Use **std::ref** to pass by **reference**

Function
parameters

```
std::thread t1(AccumulateRange,
```

```
std::ref(partial_sums[0]), 0, 1000/2 );
```



Create AND
Start each
thread
Wait for
threads to
end

```
int main() {
    const int number_of_threads = 2;
    const int number_of_elements = 1000 * 1000 * 1000;
    const int step = number_of_elements / number_of_threads;
    std::vector<uint64_t> partial_sums(number_of_threads);

    std::thread t1(AccumulateRange, std::ref(partial_sums[0]), 0, step);
    std::thread t2(AccumulateRange, std::ref(partial_sums[1]), step,
                  number_of_threads * step);

    t1.join();
    t2.join();

    uint64_t total =
        std::accumulate(partial_sums.begin(), partial_sums.end(), uint64_t(0));
    PrintVector(partial_sums); // Prints partial_sums
    std::cout << "total: " << total << std::endl;

    return 0;
}
```



Vector of Threads



Vector of
threads
Create, start,
and push each
thread into a
vector
Wait for
threads to
end

```
int main() {
    const int number_of_threads = 10;
    uint64_t number_of_elements = 1000 * 1000 * 1000;
    uint64_t step = number_of_elements / number_of_threads;
    std::vector<std::thread> threads;
    std::vector<uint64_t> partial_sums(number_of_threads);
    for (uint64_t i = 0; i < number_of_threads; i++) {
        threads.push_back(std::thread(AccumulateRange, std::ref(partial_sums[i]),
                                      i * step, (i + 1) * step));
    }
    for (std::thread &t : threads) {
        if (t.joinable()) {
            t.join();
        }
    }
    uint64_t total =
        std::accumulate(partial_sums.begin(), partial_sums.end(), uint64_t(0));
    PrintVector(partial_sums);
    std::cout << "total: " << total << std::endl;

    return 0;
}
```



```
std::thread t1(AccumulateRange,  
               std::ref(partial_sums[0]), 0, 1000/2 );
```

What do you need to take away?

- `std::thread` creates a **new** thread.
 - The **first parameter**: the name of the **function**.
 - The **rest of the parameters** will be passed to `function`.
 - All parameters passed to `function` are **passed by value**.
 - For pass by reference, wrap them in **`std::ref`**.
- **No** return value!
 - Store return value in one of the parameters passed by reference.
- Each thread **starts** as soon as it gets **created**.
- We use **`join()`** function to wait for a thread to finish



Create Threads Using Functors



Reminder: a functor is a **class/struct** that defines **operator()**

```
// comparator predicate: returns true if a < b, false otherwise
struct IntComparator
{
    bool operator()(const int &a, const int &b) const
    {
        return a < b;
    }
};

int main()
{
    std::vector<int> items { 4, 3, 1, 2 };
    std::sort(items.begin(), items.end(), IntComparator());
    return 0;
}
```

Wikipedia



Calculating the sum of numbers in the range [**start**, **end**)

```
class AccumulateFunctor {
public:
    void operator()(uint64_t start, uint64_t end) {
        _sum = 0;
        for (auto i = start; i < end; i++) {
            _sum += i;
        }
        std::cout << _sum << std::endl;
    }
    uint64_t _sum;
};
```



Vector of
threads
Vector of
functors
Create, start,
and push each
thread into a
vector
Wait for
threads to
end

Calculate
total sum

```
const int number_of_threads = 10;
uint64_t number_of_elements = 1000 * 1000 * 1000;
uint64_t step = number_of_elements / number_of_threads;
std::vector<std::thread> threads;
std::vector<AccumulateFunctor *> functors;
for (int i = 0; i < number_of_threads; i++) {
    AccumulateFunctor *functor = new AccumulateFunctor();
    threads.push_back(
        std::thread(std::ref(*functor), i * step, (i + 1) * step));
    functors.push_back(functor);
}
for (std::thread &t : threads) {
    if (t.joinable()) {
        t.join();
    }
}
int64_t total = 0;
for (auto pf : functors) {
    total += pf->_sum;
}
std::cout << "total: " << total << std::endl;
```



```
AccumulateFunctor *functor = new AccumulateFunctor();  
std::thread(std::ref(*functor), 0, 1000/2);
```

What do you need to take away?

- Creates the first parameter is either:
 - The **functor object**
 - If you don't need to use member variables later
 - a **reference to the functor object**.
 - If you need to store in and use the member variables later
- The rest of the parameters are passed to the **operator()** function.
- Return value can be **stored** in a **member variable**
 - As an alternative to passing a reference



Create Threads Using Lambda Functions



Reminder: Lambda function is a function definition that is not bound to an identifier.

```
[capture](parameters) -> return_type { function_body }
```

```
[](int x, int y) -> int { return x + y; }
```

```
std::vector<int> some_list{ 1, 2, 3, 4, 5 };  
int total = 0;  
std::for_each(begin(some_list), end(some_list),  
              [&total](int x) { total += x; });
```

Wikiped
ia



Calculating the sum of numbers in the range [**start**, **end**)

```
[i, &partial_sums, step] {  
    for (uint64_t j = i * step; j < (i + 1) * step; j++) {  
        partial_sums[i] += j;  
    }  
}
```

```
for (uint64_t i = 0; i < number_of_threads; i++) {  
    threads.push_back(std::thread([i, &partial_sums, step] {  
        for (uint64_t j = i * step; j < (i + 1) * step; j++) {  
            partial_sums[i] += j;  
        }  
    }));  
}
```



Vector of
threads
Vector of sums

Create, start,
and push each
thread into a
vector

Wait for
threads to
end
Calculate
total sum

```
const int number_of_threads = 10;
uint64_t number_of_elements = 1000 * 1000 * 1000;
uint64_t step = number_of_elements / number_of_threads;
std::vector<std::thread> threads;
std::vector<uint64_t> partial_sums(number_of_threads);
for (uint64_t i = 0; i < number_of_threads; i++) {
    threads.push_back(std::thread([i, &partial_sums, step] {
        for (uint64_t j = i * step; j < (i + 1) * step; j++) {
            partial_sums[i] += j;
        }
    }));
}
for (std::thread &t : threads) {
    if (t.joinable()) {
        t.join();
    }
}
uint64_t total =
    std::accumulate(partial_sums.begin(), partial_sums.end(), uint64_t(0));
PrintVector(partial_sums);
std::cout << "total: " << total << std::endl;
```



```
std::thread([i, &partial_sums, step] {  
    for (uint64_t j = i * step; j < (i + 1) * step; j++) {  
        partial_sums[i] += j;  
    }  
});
```

What do you need to take away?

- As an alternative to passing a parameter, we can pass references to lambda functions using **lambda capture**.



std::async

Task, Futures, and Promises



Pointer to Function

```
void AccumulateRange(uint64_t &sum, uint64_t start,
                    uint64_t end) {
    sum = 0;
    for (uint64_t i = start; i < end; i++) {
        sum += i;
    }
}
```

Function

```
std::thread t1(AccumulateRange,
```

```
std::ref(partial_sums[0]), 0, 1000/2 );
```

```
t1.join();
```

Function
parameters



Pointer to Function

Function

```
uint64_t GetRangeSum(uint64_t start, uint64_t end) {  
    uint64_t sum = 0;  
    for (uint64_t i = start; i < end; i++) {  
        sum += i;  
    }  
    return sum;  
}
```

Return value

```
auto t = std::async(GetRangeSum, 0, 100/2)
```

Future

```
return_value = t.get()
```

Function
parameters



Vector of tasks

Create, start,
and push each
task into a
vector

Wait for
tasks to end
and read
return values

```
const int number_of_threads = 10;
uint64_t number_of_elements = 1000 * 1000 * 1000;
uint64_t step = number_of_elements / number_of_threads;
std::vector<std::future<uint64_t>> tasks;

for (uint64_t i = 0; i < number_of_threads; i++) {
    tasks.push_back(std::async(GetRangeSum, i * step, (i + 1) * step));
}

uint64_t total = 0;
for (auto &t : tasks) {
    total += t.get();
}

std::cout << "total: " << total << std::endl;
```



- The **std::async** can take one more parameter of type **std::launch**. It can take the following values:
 - **std::launch::async**: the function will run on its own (new) thread.
 - **std::launch::deferred**: that the function call will be deferred until either `wait()` or `get()` is called on the future.
 - **std::launch::async | std::launch::deferred**: that the implementation may choose. (**Default option**)

```
auto t = std::async(policy, GetRangeSum, 0, 100/2)
```



How many threads?

```
unsigned int n =  
std::thread::hardware_concurrency();  
std::cout << n << " concurrent threads are  
supported.\n";
```

Returns the number of concurrent threads supported by the implementation. The value should be considered only a hint.



```
auto t = std::async(GetRangeSum, 0, 100/2)
```

What do you need to take away?

- **Tasks** are created using **std::async**.
- **Future**: The **returned value** from **std::async**
 - We get Future's value by calling **get()**
- If the future values are not ready, the main thread **blocks** until the future value becomes ready (similar to **join()**).
- **Promise**: Return value of the function passed to **std::async**.
 - For the most part, you don't need to know details of **std::promise** or define any variable of type **std::promise**. The C++ library does that behind the scenes.
- Tasks start based on the given **std::launch** policy.



Summary of Threads/Async



Summary

- **Threads**: Create using `std::thread`
 - Doesn't return a value
- **Tasks**: Create using `std::async`
 - Returns a value
- Both can use:
 - Pointer to function
 - Functor
 - Lambda functions

