




Introduction to Compilers, Part 3

ITP 435
Week 11, Lecture 1


Three Stage "Optimizing" Compiler







- Given an IR, we must generate assembly that will run on the target machine in question
- Our Expr tree:



Our ITP-11 Architecture



- It's very RISC-like (kinda MIPSy)
- Many ops have a destination and two sources:
add r1,r2,r3
- (Commonly called three address code)

Adding Code Gen to the AST Node (Simple Case)



```
struct CodeContext
{
    std::vector<std::string> ops;
};

class Node
{
public:
    void codeGen(CodeContext& c) = 0;
};
```

Recommended CodeContext starting point



```
struct CodeContext
{
    // Store the op name separate from operands
    std::vector<std::pair<std::string,
                    std::vector<strings>>> ops;
};

class Node
{
public:
    void codeGen(CodeContext& c) = 0;
};
```

Adding Code Gen to the AST Node, Cont'd



```
void NAdd::codegen( CodeContext& c )
{
    mLhs->codegen(c);
    mRhs->codegen(c);

    // This is specifically from calc example
    // you won't have this function
    AddOp(c, std::string("add"));
}
```




Register Allocation

Problem: Register Allocation



- We have a limited number of registers
- We want data to be in registers as much as possible, because it's way more efficient
- Given the following expression:


$$5+12/(4+2)$$

What is the most efficient way to calculate this in x86 assembly?


Register Allocation – What's the Big Deal?



Memory Caching



AST for $5+12/(4+2)$



Post-Order Traversal: 5,12,4,2,+,/,+

Possible Solution?



- Always load an integer in a register
- Perform each operation as it comes up in the post-order traversal on the registers assigned to the lhs/rhs
- ??
- Profit!

In an ideal world...



5	eax <-- 5
12	ebx <-- 12
4	ecx <-- 4
2	edx <-- 2
+	ecx <-- ecx + edx
/	ebx <-- ebx / ecx
+	eax <-- eax + ebx

Problem: 4 Registers May Not Be Enough



- An expression that is a very imbalanced AST will need more than 4 registers...
- For instance:
- $5 + 12 / (4 + 2 * (5 + 3))$
- This needs a minimum of **5** registers if we use that previous scheme.

Problem 2: Multiply/Divide in x86



- Addition/subtraction can be done with arbitrary registers.
- This is essentially `eax += ebx`:
`add eax, ebx`
- **Multiplication/division cannot be done with arbitrary registers.**
- Multiplication is between `eax` and the register you supply
- Multiplication results in a 64-bit number. The 32 least significant bits are stored in `eax`, the 32 most significant bits are stored in `edx`
- Eg. this is essentially `edx:eax = eax * ebx`:
`imul ebx`

Solution for Calc 3 Example...



- Integers are pushed onto the CPU stack
- When it's time to do an operation, values are popped off the stack, the result is computed and pushed back onto the stack
- This is essentially the way calculation was done in the first calc example, but now with assembly magics!



- Here is the x86 assembly generated for $5+12/(4+2)$

```
push 5
push 12
push 4
push 2
pop ebx
pop eax
add eax,ebx
push eax
pop ebx
pop eax
cdq
idiv ebx
push eax
pop ebx
pop eax
add eax,ebx
push eax
```



- Let's take a look at calc 3 now...

5+12/(4+2) Assembly, Second Look



- The assembly calc 3 generated is actually pretty terrible:

```
push 5
push 12
push 4
push 2
pop ebx
pop eax
add eax,ebx
push eax
pop ebx
pop eax
cdq
idiv ebx
push eax
pop ebx
pop eax
add eax,ebx
push eax
```

Hand-Written x86

```
mov ebx,4
mov ecx,2
add ebx,ecx
mov eax,12
cdq
idiv ebx
mov ebx,5
add eax,ebx
```



Virtual Registers




- Don't assign the actual registers during the CodeGen pass, instead use *virtual* registers
- Assume we have an infinite number of virtual registers
- This makes it easier to code gen without worrying about register limitations

AST for $5+12/(4+2)$




- During the CodeGen post-order traversal, we assign virtual registers as the nodes are generated



AST for $5+12/(4+2)$




- During the CodeGen post-order traversal, we assign virtual registers as the nodes are generated
- Denoted by % in front of the number (just to distinguish it from a real register)



In ITP-11 Code...



```
movi %0,5  
movi %1,12  
movi %2,4  
movi %3,2  
add %4,%2,%3  
div %5,%1,%4  
add %6,%0,%5
```



Register Allocation Process



- When we do register allocation (in a bit), we will then need to convert virtual registers into actual registers on the target machine
- For ITP-11, we want to only allocate r1 through r7, since the other registers are used for special cases

```
movi %0,5  
movi %1,12  
movi %2,4  
movi %3,2  
add %4,%2,%3  
div %5,%1,%4  
add %6,%0,%5
```



CodeGen for Arrays

The Problem



- Suppose I have the following statements:

`abc = 20;`

`a = 2`

`myarr[a] = abc * 5;`

- *How do I calculate the correct stack offset for this array index?*

Array Memory Address Calculation



- Given:
 - base = Base address where the array starts at (this is a known constant in our case)
 - subscript = subscript index you want (this could be a variable)
 - size = sizeof each element of type (in PA6 this is 1)

$$\text{address} = \text{base} + \text{subscript} * \text{size}$$

OR in PA6...

$$\text{address} = \text{base} + \text{subscript}$$

What's the subscript?



```
abc = 20;
```

```
a = 2
```

```
myarr[a] = abc * 5;
```

1. Code gen the expression inside of the []. This will give you a subscript in a temp virtual register
2. movi the base address into another temp virtual register
3. Generate an add *instruction* between 1 and 2 to get the actual stack offset you want!



CodeGen for If and While




- Jumps need to have the instruction number specified (cause they set `pc = location`)

```
if abc == 1 {  
    abc = 15;  
    xyz = 20;  
    // other stuff...  
}
```

- When we generate the condition, we don't know the target of the jump until we generate the block inside the if

The Problem



- Each statement may be several instructions long. We don't know how many until we generate the block

Solution



- Suppose we so far have 10 other instructions prior to the if
- Next instruction will be at index 10, then
- (We'll just assume there's no prior virtual registers used)

#	Instruction
9	...



- First, generate the condition, which:
- First generates the left hand side

```
if abc == 1 {  
    abc = 15;  
    xyz = 20;  
}
```

- {abc} is the constant stack index we're using for abc, which we'd know cause we saved it in the CodeContext

#	Instruction
9	...
10	loadi %0,{abc}

Solution



- Next, generate the expression for the right hand side of the condition

```
if abc == 1 {  
    abc = 15;  
    xyz = 20;  
}
```

#	Instruction
9	...
10	loadi %0,{abc}
11	movi %1,1

Solution



- Next, generate the comparison

```
if abc == 1 {  
    abc = 15;  
    xyz = 20;  
}
```

#	Instruction
9	...
10	loadi %0,{abc}
11	movi %1,1
12	cmpeq %0,%1

Solution



- Now generate the “jump if not true”, with a placeholder address

```
if abc == 1 {  
    abc = 15;  
    xyz = 20;  
}
```

- Since we don't know the address for the end of the if block yet, we leave the number for the movi blank (denoted by ??)

#	Instruction
9	...
10	loadi %0,{abc}
11	movi %1,1
12	cmpeq %0,%1
13	movi %2,??
14	jnt %2

Solution



- Now generate the block of statements (which would each individually CodeGen...)

```
if abc == 1 {  
    abc = 15;  
    xyz = 20;  
}
```

#	Instruction
9	...
10	loadi %0,{abc}
11	movi %1,1
12	cmpeq %0,%1
13	movi %2,??
14	jnt %2
15	movi %3,15
16	storei {abc},%3
17	movi %4,20
18	storei {xyz},%4

Solution



- Once we're done with the block, we know that the "if" block ends at instruction 18...
- So we should fix the movi to be one past the end, or 19!

#	Instruction
9	...
10	loadi %0,{abc}
11	movi %1,1
12	cmpeq %0,%1
13	movi %2,19
14	jnt %2
15	movi %3,15
16	storei {abc},%3
17	movi %4,20
18	storei {xyz},%4



- Same idea, just a little more gnarly:
 1. Generate if condition
 2. Generate jnt w/ placeholder address (to else block)
 3. Generate if block code
 4. Generate unconditional jmp w/ placeholder address (to past else block)
 5. Fix up address in #2
 6. Generate else block code
 7. Fix up address in #4

```
if abc < 37 {  
    x = 1;  
} else {  
    x = 0;  
}
```



- Similar as well:
 1. Generate while condition
 2. Generate jnt w/ placeholder address (to past while loop)
 3. Generate while block code
 4. Generate unconditional jmp back to first instruction of #1
 5. Fix up address in #2

```
while x < count {  
    myarr[x] = x * 5;  
    ++x;  
}
```



Linear Scan Register Allocator

Register Allocation Process



- We need to convert virtual registers into actual registers on the target machine
- For ITP-11, we want to only allocate r1 through r7, since the other registers are used for special cases

```
movi %0,5  
movi %1,12  
movi %2,4  
movi %3,2  
add %4,%2,%3  
div %5,%1,%4  
add %6,%0,%5
```

Methods for Register Allocation



- Register allocation is an NP-hard problem
- There are several methods to allocate registers, including graph coloring
- We're going to use a fairly simple approach called *linear scan*

Linear Scan: Code w/ Virtual Registers



- Suppose our code gen yielded the following sequence of instructions using virtual registers

#	Instruction
0	push 0
1	movi %0,5
2	movi %1,12
3	movi %2,4
4	movi %3,2
5	add %4,%2,%3
6	div %5,%1,%4
7	add %6,%0,%5
8	storei 0, %6

Linear Scan: Generate Intervals



- First, we need to generate the intervals for each virtual register, which tells us the first and last instruction a VR is used at

#	Instruction
0	push 0
1	movi %0,5
2	movi %1,12
3	movi %2,4
4	movi %3,2
5	add %4,%2,%3
6	div %5,%1,%4
7	add %6,%0,%5
8	storei 0, %6

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

Linear Scan: The Scan



- Before we start
 - Virtual registers have no real register assignments
 - All real registers are available

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	
%1	
%2	
%3	
%4	
%5	
%6	

R	Available?
r1	true
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true



- Next we loop from 0 to # of instructions, and at each instruction:
 - Decide if any VR just expired (we're at last for that VR) and if so, whichever real register it was assigned to is now free
 - Decide if any VR now is active (we're at first for that VR) and if so, assign first available real register

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

Linear Scan: The Scan



- Instruction 0
 - Nothing is expired
 - Nothing is active yet

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	
%1	
%2	
%3	
%4	
%5	
%6	

R	Available?
r1	true
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 1
 - Nothing is expired

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	
%1	
%2	
%3	
%4	
%5	
%6	

R	Available?
r1	true
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 1
 - Nothing is expired
 - %0 is now active, so assign first available (r1)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	
%2	
%3	
%4	
%5	
%6	

R	Available?
r1	false
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 2
 - Nothing is expired

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	
%2	
%3	
%4	
%5	
%6	

R	Available?
r1	false
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true



Linear Scan: The Scan

- Instruction 2
 - Nothing is expired
 - %1 is now active, so assign first available (r2)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	
%3	
%4	
%5	
%6	

R	Available?
r1	false
r2	false
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 3
 - Nothing is expired
 - %2 is now active, so assign first available (r3)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	
%4	
%5	
%6	

R	Available?
r1	false
r2	false
r3	false
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 4
 - Nothing is expired
 - %3 is now active, so assign first available (r4)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	
%5	
%6	

R	Available?
r1	false
r2	false
r3	false
r4	false
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 5
 - Both %2 and %3 expired, so set their real registers to available

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	
%5	
%6	

R	Available?
r1	false
r2	false
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 5
 - Both %2 and %3 expired, so set their real registers to available
 - %4 is now active, so assign first available (r3)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	
%6	

R	Available?
r1	false
r2	false
r3	false
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 6
 - %1 and %4 are now expired, so their real registers are available

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	
%6	

R	Available?
r1	false
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 6
 - %1 and %4 are now expired, so their real registers are available
 - %5 is now active, so assign first available (r2)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	

R	Available?
r1	false
r2	false
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 7
 - %0 and %5 are now expired, so r1/r2 are available

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	

R	Available?
r1	true
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true



Linear Scan: The Scan

- Instruction 7
 - %0 and %5 are now expired, so r1/r2 are available
 - %6 is now active, so assign first available (r1)

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	r1

R	Available?
r1	false
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 8
 - %6 is now expired, so r1 is available

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	r1

R	Available?
r1	true
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: The Scan



- Instruction 8
 - %6 is now expired, so r1 is available
 - Nothing new is active

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	r1

R	Available?
r1	true
r2	true
r3	true
r4	true
r5	true
r6	true
r7	true

Linear Scan: Result



- Given the code with virtual registers and the allocations, we simply replace the virtual registers with the real ones

#	Instruction
0	push 0
1	movi %0,5
2	movi %1,12
3	movi %2,4
4	movi %3,2
5	add %4,%2,%3
6	div %5,%1,%4
7	add %6,%0,%5
8	storei 0, %6

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	r1

#	Instruction
0	push 0
1	movi r1,5
2	movi r2,12
3	movi r3,4
4	movi r4,2
5	add r3,r3,r4
6	div r2,r2,r3
7	add r1,r1,r2
8	storei 0, r1

Linear Scan: The Scan (Summary)



- Loop from 0 to # of instructions, and at each iteration we:
 - Decide if any VR just expired (we're at last for that VR) and if so, whichever real register it was assigned to is now free
 - Decide if any VR now is active (we're at first for that VR) and if so, assign first available real register

VR	First	Last
%0	1	7
%1	2	6
%2	3	5
%3	4	5
%4	5	6
%5	6	7
%6	7	8

VR	Real
%0	r1
%1	r2
%2	r3
%3	r4
%4	r3
%5	r2
%6	r1

Example from PA6: High-level Code



```
data {  
    var x;  
}  
main {  
    x = 10 + 5;  
    x = x * 2;  
    x = (x + 20) / 3;  
}
```



```
0: push r0
1: movi %0,10
2: movi %1,5
3: add %2,%0,%1
4: storei 0,%2
5: loadi %3,0
6: movi %4,2
7: mul %5,%3,%4
8: storei 0,%5
9: loadi %6,0
10: movi %7,20
11: add %8,%6,%7
12: movi %9,3
13: div %10,%8,%9
14: storei 0,%10
15: exit
```

What are the intervals of the virtual registers?

Example from PA6



	INTERVALS:	ALLOCATION:	
push r0			push r0
movi %0,10	%0:1,3	%0:r1	movi r1,10
movi %1,5	%1:2,3	%1:r2	movi r2,5
add %2,%0,%1	%2:3,4	%2:r1	add r1,r1,r2
storei 0,%2	%3:5,7	%3:r1	storei 0,r1
loadi %3,0	%4:6,7	%4:r2	loadi r1,0
movi %4,2	%5:7,8	%5:r1	movi r2,2
mul %5,%3,%4	%6:9,11	%6:r1	mul r1,r1,r2
storei 0,%5	%7:10,11	%7:r2	storei 0,r1
loadi %6,0	%8:11,13	%8:r1	loadi r1,0
movi %7,20	%9:12,13	%9:r2	movi r2,20
add %8,%6,%7	%10:13,14	%10:r1	add r1,r1,r2
movi %9,3			movi r2,3
div %10,%8,%9			div r1,r1,r2
storei 0,%10			storei 0,r1
exit			exit



- For a real program, we still might run out of registers
- Ultimately some stuff needs to be “spilled” into memory
- Don’t have to worry about this in our case


Graph Coloring Allocator




Code Sequence

```
A = ...  
B = ...  
... B ...  
C = ...  
... A ...  
D = ...  
... D ...  
... C ...
```

Interference Graph



Colored Graph



Final Allocation

```
R1 = ...  
R2 = ...  
... R2 ...  
R2 = ...  
... R1 ...  
R1 = ...  
... R1 ...  
... R2 ...
```

- We won't do this, but another popular method
- Convert intervals into interference graph and then colors graph
- Gives better allocations than linear scan, though more expensive to compute